


Full-design latency reduction vs. each reference



Reduction = (reference - full design) / reference.
FP16 evaluates the overall quantization framework.