



(a) Group quantization

W8A8 G64 | W4A8 G32

A [M, K]

B [K, N]




(b) Scale-aligned merge




$$c_g = \text{RNE}(s_g \cdot 2^{24-E})$$

Scale coefficients



(c) Output paths

FP16 main path



Optional INT8 path

